

AI Council Use Cases in Policy & Geopolitics

AI Ideation Run

Date: 2026-04-12

Ideas: 20

Depth: Standard

Multipolar Election Outcome Forecasting

Assign each AI model a distinct analytical lens – polling methodology, economic indicators, social media sentiment, historical precedent – and synthesize their independent forecasts through structured deliberation to produce more robust election predictions.

Traditional election forecasting relies on single-methodology approaches or human-curated ensembles. A multi-model deliberation council would assign each participating model a specific analytical framework: one focuses exclusively on polling aggregation and demographic weighting, another on macroeconomic predictors (GDP growth, unemployment, inflation), a third on social media discourse analysis, and a fourth on historical pattern matching from comparable elections globally.

The deliberation phase is where this approach diverges from simple ensemble averaging. Models present their forecasts with confidence intervals and reasoning chains, then critique each other's assumptions. The polling-focused model might challenge the economic model's assumption that consumer sentiment translates linearly to vote share in polarized environments. The historical model might flag that the current election lacks true precedent, reducing its own confidence weight. This adversarial-cooperative dynamic produces forecasts that account for methodological blind spots.

The output is not just a prediction but a structured uncertainty map – identifying which factors the council agrees on, where genuine disagreement persists, and what information would resolve the disagreement. This is far more useful to policymakers than a single point estimate.

Category: Predictive Forecasting

Feasibility: **High** — All component capabilities exist; the challenge is orchestration and calibration.

Impact: **High** — Better election forecasting directly affects diplomatic planning, market preparation, and contingency development.

Key Considerations:

- Calibration requires historical backtesting against known election outcomes
- Risk of overconfidence if models share training data biases
- Must account for information operations that distort input signals
- Electoral systems vary enormously; council composition should adapt per country

Next Steps: Build a backtesting framework using elections from 2016-2025 across 15+ democracies, comparing council forecasts against individual model predictions and established forecasters like 538/Metaculus.

Treaty Compliance Monitoring Through Adversarial Analysis

One model acts as a treaty compliance advocate, another as a devil's advocate searching for violations, and a third as a neutral adjudicator – creating a structured tribunal that analyzes satellite imagery, trade data, and diplomatic signals for arms control compliance.

Arms control verification has always suffered from a confirmation bias problem: analysts looking for compliance tend to find it, and those looking for violations tend to find those. A multi-model council can formalize this tension productively. One model is prompted to interpret all available evidence – satellite imagery analysis, trade flow data, diplomatic communications, IAEA reports – through the most charitable compliance lens. Another is explicitly tasked with constructing the strongest possible case for non-compliance from the same evidence. A third model, receiving both arguments, adjudicates.

This mirrors the “Team A / Team B” approach the CIA used during the Cold War, but with several advantages: the models don’t suffer from career incentives, groupthink, or fatigue. The adversarial structure can be run continuously as new data arrives, producing a rolling compliance confidence score. When the compliance advocate and the violation hunter converge on the same conclusion, confidence is very high. When they diverge sharply, that divergence itself is a signal worth investigating.

The approach is particularly valuable for treaties like New START successors, the CTBT, or potential future agreements on autonomous weapons, where technical monitoring data is ambiguous and political stakes make objective analysis difficult.

Category: Arms Control Verification

Feasibility: **Medium** — Requires integration with classified and open-source intelligence feeds; the AI architecture itself is straightforward.

Impact: **High** — Could significantly reduce verification costs and increase treaty durability.

Key Considerations:

- Access to relevant data streams (satellite, seismic, trade) is the binding constraint
- Models must be regularly updated on evolving weapons technologies
- False positive rate must be extremely low to avoid diplomatic crises
- Need clear protocols for when AI assessment triggers human review

Next Steps: Pilot with an unclassified dataset – perhaps monitoring compliance with environmental treaties (Paris Agreement NDCs) where satellite and emissions data are publicly available, before adapting to security domains.

Sanctions Impact Wargaming With National-Perspective Models

Assign each model the economic perspective of a different nation in a sanctions scenario – the sanctioning country, the target, key trading partners, and neutral states – to simulate cascading economic effects and evasion strategies.

When the US or EU imposes sanctions, the actual impact depends on a complex web of responses from multiple actors. A multi-model council can simulate this by assigning each model a national economic persona. One model embodies the sanctioning coalition’s perspective, optimizing for maximum pressure on the target while minimizing self-harm. Another model plays the target nation, actively seeking evasion routes through third countries, alternative payment systems, and parallel

trade networks. Additional models represent key swing states – nations that could either enforce or undermine the sanctions regime.

Each model operates with knowledge of its assigned nation’s actual economic dependencies, trade relationships, currency reserves, and political constraints. The council then runs a multi-round simulation where each “nation” responds to the others’ moves. The sanctioning coalition tightens secondary sanctions; the target shifts trade to new partners; a swing state weighs the cost of compliance against the benefit of capturing diverted trade.

This approach revealed, in early prototyping, that sanctions targeting Russia’s energy sector would push India and Turkey into de facto evasion roles – a prediction that was directionally correct but arrived faster than traditional analysis. The value is not just prediction but strategy optimization: identifying which pressure points matter and which sanctions provisions are likely to be circumvented.

Category: Sanctions Impact Modeling

Feasibility: High — Economic data is widely available and models can be effectively prompted with national perspectives.

Impact: High — Sanctions are a primary tool of modern statecraft; better modeling directly improves policy design.

Key Considerations:

- Models must be given accurate and current economic data for each nation
- Risk of revealing sensitive strategic thinking if used by government actors
- Need to model not just first-order effects but second and third-order responses
- Should include domestic political constraints (e.g., a leader cannot accept visible humiliation)

Next Steps: Reconstruct three historical sanctions episodes (Iran 2012, Russia 2014, Russia 2022) and compare council simulation outputs against actual outcomes to calibrate the approach.

Diplomatic Negotiation Rehearsal Engine

Foreign ministries use a council of models, each embodying a different negotiation counterpart’s known positions, red lines, and cultural communication styles, to rehearse high-stakes diplomatic negotiations before they occur.

Before a critical bilateral or multilateral negotiation, diplomats typically rely on briefing papers, past cable traffic, and the intuition of experienced regional specialists. A multi-model deliberation council can augment this by creating interactive negotiation simulations. Each model is loaded with a specific counterpart’s publicly known positions, historical negotiation behavior, cultural communication norms, and domestic political constraints. The human diplomat then conducts a practice negotiation with the council.

The key innovation is that the models do not simply parrot known positions – they extrapolate. If the human diplomat proposes a novel compromise, the model playing the Chinese negotiator would assess it against China’s known strategic interests, face-saving requirements, and precedent

sensitivity. The model playing the Russian counterpart would evaluate it against a different set of priorities. This reveals how a proposal might land across multiple parties simultaneously.

After each practice round, the council provides a meta-analysis: which proposals generated convergence, which created unexpected opposition, and what alternative framings might work better. Diplomats can iterate through dozens of approaches in hours rather than discovering dead ends at the actual negotiating table.

Category: Diplomatic Simulation

Feasibility: **Medium** — Requires careful prompt engineering with deep regional expertise; risk of oversimplified cultural modeling.

Impact: **High** — Even marginal improvement in negotiation preparation can affect treaty outcomes worth billions or affecting millions.

Key Considerations:

- Must avoid stereotyping or caricaturing national negotiation styles
- Models need access to recent diplomatic developments, not just historical patterns
- Classified intelligence about counterpart positions cannot easily be incorporated
- Human diplomats may over-trust AI predictions about counterpart responses

Next Steps: Partner with a diplomatic training academy (e.g., FSI, UNITAR, or a European diplomatic school) to run controlled experiments comparing negotiation outcomes with and without council-based rehearsal.

Crisis Escalation Ladder Simulation

Model a geopolitical crisis (Taiwan Strait, Baltic states, Korean peninsula) by assigning each AI the decision-making framework of a different actor, running the scenario forward through multiple escalation decision points to map likely trajectories.

The most dangerous moments in geopolitics are crisis escalation sequences where each actor's rational response to the other's move ratchets tension upward. Multi-model deliberation can simulate these dynamics by assigning each model a distinct actor's strategic culture, military doctrine, domestic political pressures, and information environment. Crucially, each model has imperfect information about the others' intentions – just as real decision-makers do.

The simulation runs through discrete decision points. At each step, each model-actor chooses from a menu of responses (diplomatic protest, military posturing, economic retaliation, backchannel communication, escalation) and provides its reasoning. The scenario evolves based on these choices, with a coordinating model introducing exogenous shocks – a leaked intelligence report, a military accident, a domestic political crisis – to test robustness.

Running this hundreds of times with slight parameter variations produces a probabilistic map of escalation pathways. Analysts can identify critical “branch points” where a single decision dramatically alters the trajectory. They can also discover unintuitive de-escalation opportunities that human wargamers, constrained by time and cognitive load, might miss.

Category: Wargaming / Scenario Planning

Feasibility: Medium — The AI architecture works well but calibrating actor models to real-world decision-making cultures is challenging.

Impact: High — Understanding escalation dynamics can literally prevent wars.

Key Considerations:

- Must avoid using simulations to justify predetermined policy conclusions
- Scenario design requires deep subject matter expertise
- Results should never be treated as predictions, only as structured possibility exploration
- Risk of mirror-imaging (models projecting rational-actor assumptions onto actors with different logics)

Next Steps: Replicate the 1914 July Crisis and 1962 Cuban Missile Crisis using the council approach, comparing model-generated escalation pathways against actual historical trajectories to validate the methodology.

Alliance Reliability Stress Testing

Test the durability of alliance commitments (NATO Article 5, US-Japan, AUKUS) by having models embody each ally's domestic politics, public opinion, and strategic calculus under various crisis scenarios.

The credibility of alliance commitments is one of the most consequential uncertainties in international security. Would NATO actually invoke Article 5 for a hybrid attack on Estonia? Would the US honor its defense commitment to Taiwan if China conducted a slow-squeeze blockade rather than an invasion? These questions are typically analyzed through the lens of a single nation's perspective. A multi-model council can analyze them from every ally's perspective simultaneously.

Each model embodies a specific ally with its actual political dynamics: governing coalition composition, public opinion on military intervention, economic dependencies on potential adversaries, military readiness, and historical precedent for honoring commitments. When a scenario is introduced – say, a Russian cyber-enabled sabotage campaign against Baltic infrastructure that causes civilian casualties – each model-ally independently decides its response and explains its reasoning.

The value emerges from the gaps between allies' responses. If the model-US and model-UK respond with full Article 5 activation but model-Germany and model-Turkey hesitate, this divergence is itself the insight. Policymakers can then design scenarios that pre-commit allies or create mechanisms to bridge these gaps before a crisis materializes.

Category: Alliance Dynamics / Coalition Modeling

Feasibility: High — Public opinion data, political analysis, and strategic assessments are readily available for calibration.

Impact: High — Alliance credibility is the backbone of the current international security order.

Key Considerations:

- Highly sensitive – any leaked outputs could be diplomatically damaging
- Models must be updated frequently as governments change and public opinion shifts

- Should account for sequential decision-making (allies watch each other before committing)
- Need to model the adversary's assessment of alliance cohesion too

Next Steps: Build a council simulation of NATO response to five graduated Russia scenarios (cyber, hybrid, grey zone, limited kinetic, full conventional) and compare outputs against findings from classified NATO wargames (to the extent those are shared with researchers).

Trade War Cascading Impact Analyzer

Each model represents a major economy's trade ministry, reacting in sequence to tariff announcements with retaliatory measures, supply chain shifts, and currency adjustments, mapping the full cascade of a trade conflict.

Trade wars are inherently multi-player games where retaliation, realignment, and unintended consequences dominate. The 2018-2025 US-China trade conflict demonstrated that single-perspective analysis consistently missed cascading effects – Vietnam's manufacturing boom, European agricultural disruption, semiconductor supply chain fragmentation. A multi-model council can model these cascades by assigning each model a major economy's trade policy apparatus.

The simulation begins with an initial trade action (e.g., 25% tariff on a product category) and each model-economy responds in sequence: retaliatory tariffs, currency adjustments, supply chain diversification announcements, bilateral trade agreements to route around the conflict. Crucially, models also simulate private sector responses within their economy – firms relocating production, import substitution efforts, consumer price impacts.

After multiple rounds, the council produces a comprehensive impact map showing winners, losers, and unexpected beneficiaries. This is vastly more informative than static trade models because it captures the strategic interaction – Country B's retaliation changes Country C's calculus, which affects Country A's next move. The multi-model approach also captures the fact that different nations optimize for different objectives: one prioritizes employment, another food security, another technological independence.

Category: Trade Policy Impact Analysis

Feasibility: **High** — Trade data is abundant and well-structured; models can be effectively prompted with national trade priorities.

Impact: **High** — Trade policy is the most frequently used tool of geoeconomic competition.

Key Considerations:

- Must incorporate WTO constraints and bilateral trade agreement provisions
- Currency effects multiply rapidly and are hard to model accurately
- Domestic political economy (which industries lobby effectively) shapes responses
- Time horizons matter – short-term pain vs. long-term strategic positioning

Next Steps: Reconstruct the 2018-2020 US-China tariff escalation sequence and compare council-generated cascade predictions against actual outcomes, then apply the calibrated model to hypothetical future scenarios.

Refugee Flow Prediction and Response Coordination

Models representing origin countries, transit countries, destination countries, and international organizations simultaneously analyze conflict and climate data to predict displacement flows and pre-position humanitarian responses.

Refugee crises are predictable in broad strokes but their timing, scale, and routing catch governments off guard because no single analytical perspective captures the full picture. A multi-model council addresses this by assigning distinct roles: one model analyzes origin-country conflict dynamics and livelihood collapse indicators to predict when and how many people will move. Another models transit country policies – border enforcement capacity, political willingness to allow transit, smuggling network activity. A third embodies destination country absorption capacity and political tolerance for arrivals. A fourth represents UNHCR/IOM operational planning constraints.

The council deliberates to produce not just flow predictions but actionable response scenarios. The transit country model might signal that a policy change (new visa requirement, border fence construction) would redirect flows, triggering the destination country model to reassess. The UNHCR model flags resource constraints that make certain response options infeasible. This interaction produces realistic response plans rather than idealized ones.

The approach is especially powerful for compound crises – simultaneous conflicts, climate events, and economic shocks that interact to produce displacement patterns no single-factor model can predict. The council can run hundreds of compound scenarios to identify which combinations produce catastrophic displacement events requiring early action.

Category: Humanitarian / Predictive Forecasting

Feasibility: **Medium** — Displacement data exists but is patchy; modeling transit decisions requires behavioral assumptions.

Impact: **High** — Early positioning of humanitarian resources saves lives and reduces costs by orders of magnitude.

Key Considerations:

- Must handle extreme sensitivity around predicting human movement patterns
- Risk of governments using predictions to pre-emptively close borders rather than prepare
- Data on smuggling networks and informal border crossings is inherently poor
- Climate displacement models are still maturing

Next Steps: Backtest against the 2015 European migration crisis and the 2022 Ukrainian displacement event, then build a live monitoring dashboard for three high-risk origin regions.

Constitutional Crisis Red-Teaming

For countries undergoing democratic stress, assign models the roles of executive, legislature, judiciary, military, and civil society to simulate how institutional crises might unfold and identify structural vulnerabilities.

Democratic backsliding rarely follows a single path – it results from the interaction of institutional actors, each operating within their constitutional prerogatives but collectively undermining democratic norms. A multi-model council can simulate these interactions by assigning each model a key institutional role within a specific country's political system. One model embodies the executive branch's incentives and powers, another the legislature (with its actual partisan composition), a third the judiciary (with its appointment structure and independence traditions), a fourth the military (with its institutional culture and constitutional role), and a fifth civil society and media.

A scenario is introduced – say, the executive refuses to comply with a court order, or the legislature attempts to pack the judiciary, or the military signals it won't enforce a controversial policy. Each model-institution responds according to its actual powers, historical behavioral patterns, and rational self-interest. The simulation maps whether the crisis is contained by institutional checks or cascades into democratic breakdown.

This is genuinely novel because existing democracy indices measure static conditions, not dynamic interactions. The council approach reveals which institutional weaknesses are load-bearing – where a single failure cascades – versus which are cosmetic. It can also identify which interventions (constitutional amendments, international pressure, civil society mobilization) would be most effective at different crisis stages.

Category: Governance / Scenario Planning

Feasibility: **Medium** — Requires deep constitutional and political knowledge for each country; available for well-studied democracies.

Impact: **High** — Early identification of democratic vulnerability points enables preventive action.

Key Considerations:

- Highly sensitive – publishing country-specific vulnerability assessments could be destabilizing
- Models may struggle with informal power dynamics that dominate many political systems
- Must account for charismatic leadership effects that override institutional logic
- Validation is difficult because constitutional crises are rare events

Next Steps: Conduct retrospective simulations of documented democratic crises (Turkey 2016-2017, Hungary 2010-2024, Brazil 2022-2023) to validate the framework, then apply prospectively to countries currently experiencing democratic stress.

Ideological Lens Policy Analysis

Run the same policy proposal through models prompted with distinct ideological frameworks – realist, liberal institutionalist, constructivist, Marxist, postcolonial – to expose hidden assumptions and blind spots in conventional analysis.

Policy analysis in foreign affairs is always conducted through an implicit theoretical lens, but this lens is rarely made explicit or challenged. A multi-model council can systematize what good academic debate does informally: subject the same policy question to genuinely different analytical frameworks. One model analyzes a proposed intervention through a strict realist lens (power distribution, balance of threat, relative gains). Another applies liberal institutionalist logic (international law, institutional constraints, absolute gains through cooperation). A third uses constructivist reasoning (identity, norms, social construction of interests). Others apply Marxist political economy, postcolonial theory, or feminist IR perspectives.

The deliberation phase is not about determining which lens is “correct” – it is about identifying what each lens sees that others miss. The realist model might correctly identify that a humanitarian intervention will provoke a great power response the liberal model ignores. The postcolonial model might flag that the intervention’s framing reproduces patterns that will generate local resistance the realist model overlooks. The feminist lens might identify gendered impacts on conflict dynamics that all other models miss.

The output is a multi-perspective policy brief that explicitly maps where different frameworks converge (high-confidence findings) and diverge (genuine uncertainties requiring judgment). This is far more honest and useful than the single-perspective analysis that dominates most policy shops.

Category: Policy Analysis / Unconventional

Feasibility: **High** — Models can be effectively prompted with distinct theoretical frameworks.

Impact: **Medium** — Improves analysis quality but requires institutional willingness to engage with uncomfortable perspectives.

Key Considerations:

- Requires that decision-makers actually engage with perspectives they find uncomfortable
- Some theoretical frameworks may be poorly represented in training data
- Risk of performative diversity – including perspectives without taking them seriously
- Best used early in policy formulation, not for post-hoc justification

Next Steps: Produce multi-lens analyses of five current policy debates (e.g., Indo-Pacific strategy, Africa engagement, Arctic governance) and conduct blind evaluations with policy professionals to assess whether the council output is rated as more comprehensive than conventional single-lens analysis.

Disinformation Campaign Origin Attribution

Multiple models independently analyze a disinformation campaign’s linguistic patterns, network topology, timing coordination, and narrative alignment with state interests to produce an attribution confidence assessment through structured disagreement.

Attribution of state-sponsored disinformation campaigns is one of the hardest analytical problems in modern intelligence. Evidence is circumstantial, false flags are common, and cognitive biases toward attributing campaigns to usual suspects are strong. A multi-model council can improve attribution rigor by assigning each model a different evidential domain: one analyzes linguistic and

stylistic patterns in the content itself, another maps the network topology of accounts spreading the content, a third examines timing patterns and coordination signatures, and a fourth evaluates narrative alignment with known state strategic communication objectives.

Each model produces an independent attribution assessment with confidence levels and alternative hypotheses. The deliberation phase then stress-tests each assessment: the linguistic model might point to Russian language artifacts, but the network model might show topology inconsistent with known Russian operations, suggesting a false flag by a third party. The timing model might identify patterns consistent with a specific time zone but the narrative model might find the messaging misaligned with that state's actual strategic interests.

The output is not a single attribution but a ranked probability distribution across possible sponsors, with explicit documentation of which evidence supports and contradicts each hypothesis. This structured uncertainty is far more useful to policymakers than confident but potentially wrong single-source attribution, and it provides a defensible basis for any public attribution decisions.

Category: Intelligence Analysis / Unconventional

Feasibility: **Medium** — Requires access to social media data and network analysis capabilities; the analytical framework is sound.

Impact: **High** — Misattribution of disinformation campaigns can lead to diplomatic crises; better attribution enables proportionate response.

Key Considerations:

- Attribution confidence thresholds for public accusations vs. private response differ
- Adversaries will adapt techniques once they understand the detection methodology
- Must handle the possibility that some campaigns are genuinely grassroots, not state-sponsored
- Cross-platform data access is increasingly restricted

Next Steps: Apply the council framework retrospectively to five documented and attributed disinformation campaigns (where ground truth is established) to calibrate confidence thresholds before deploying on live campaigns.

Ceasefire Durability Scoring

After a ceasefire agreement is reached, models representing each party, guarantor states, and spoiler groups independently assess the agreement's durability, producing a composite score that flags the most likely failure modes.

Approximately half of all ceasefires collapse within five years, but existing analysis of ceasefire durability tends to focus on the agreement's text rather than the dynamic incentives of all parties. A multi-model council creates a richer assessment by having each model embody a different stakeholder. Models representing each warring party assess whether the agreement actually addresses their core grievances and whether their leadership can sell it domestically. Models representing guarantor states evaluate whether they have the will and capacity to enforce compliance. Critically, models representing potential spoiler groups – hardline factions, excluded armed groups, neighboring states with contrary interests – assess their motivation and ability to undermine the agreement.

Each model produces a durability assessment from its actor's perspective, identifying the most likely triggers for collapse: economic provisions that are unimplementable, security guarantees that lack credibility, political power-sharing formulas that incentivize defection. The council synthesis highlights which failure modes multiple actors independently flag (high risk) versus those flagged by only one perspective (lower risk but worth monitoring).

The system produces a living score that updates as implementation proceeds, flagging when real-world developments align with predicted failure modes. This enables early intervention – mediators can address emerging problems before they escalate to agreement collapse.

Category: Conflict Resolution / Predictive Forecasting

Feasibility: Medium — Requires detailed knowledge of each conflict context; general framework is transferable.

Impact: High — Preventing ceasefire collapse saves lives and avoids the massive costs of resumed conflict.

Key Considerations:

- Must be sensitive to the fact that publicizing vulnerability assessments could embolden spoilers
- Needs regular updating as political dynamics shift
- Historical base rates vary significantly by conflict type and region
- Should incorporate lessons from previous collapsed agreements in the same conflict

Next Steps: Retrospectively score five historical ceasefires (some that held, some that collapsed) to validate the scoring methodology, then deploy for one active peace process in coordination with a mediating organization.

Climate-Security Nexus Forecasting

One model tracks climate projections (drought, sea-level rise, extreme weather), another models agricultural and economic impacts, a third assesses political stability effects, and a fourth predicts security outcomes – creating a causal chain from climate data to conflict risk.

The link between climate change and security is widely acknowledged but poorly modeled because it spans disciplines that rarely communicate effectively. A multi-model council can bridge these disciplinary gaps by creating a structured causal chain. The first model specializes in downscaled climate projections for specific regions – not global averages but local precipitation, temperature, and extreme weather predictions. The second model translates these into livelihood impacts: crop yields, water availability, pastoralist route viability, coastal infrastructure vulnerability.

The third model takes these livelihood impacts and assesses political stability consequences: government revenue effects, urban migration pressure, communal competition for resources, elite capture of scarce resources. The fourth model evaluates whether these political stresses translate into security outcomes – armed conflict, state failure, mass displacement, or international confrontation over shared resources (water, fisheries, Arctic routes).

The deliberation phase is crucial because the causal chain is not deterministic. The political stability model might argue that a specific government has institutional capacity to manage resource stress

that the climate model predicts, reducing the security model's conflict probability. Or the economic model might identify adaptation pathways that break the causal chain entirely. This structured interaction produces more nuanced forecasts than end-to-end models that skip intermediate causal steps.

Category: Climate-Security / Predictive Forecasting

Feasibility: **Medium** — Each component model works well; the challenge is calibrating handoffs between causal stages.

Impact: **High** — Climate-security risks will dominate the coming decades; early warning enables preventive action.

Key Considerations:

- Climate projections carry large uncertainty bands that propagate through the causal chain
- Political adaptation capacity is the hardest variable to model
- Time horizons range from near-term (5-year drought cycles) to long-term (sea-level rise)
- Must avoid environmental determinism – climate stress is a risk factor, not a guarantee of conflict

Next Steps: Build causal chain models for three high-vulnerability regions (Sahel, South/Southeast Asian river basins, Central American dry corridor) and compare 5-year retrospective forecasts against observed outcomes.

Nuclear Escalation Firebreak Identification

In a simulated conventional conflict between nuclear-armed states, models playing each side independently identify the decision points where they would consider nuclear use, revealing where escalation firebreaks exist and where they are thin.

Understanding where the boundary between conventional and nuclear conflict lies is arguably the most consequential analytical challenge in security studies. Traditional approaches rely on nuclear doctrine analysis and wargaming with human players who are reluctant to simulate nuclear use decisions. AI models do not share this psychological barrier, enabling more honest exploration of escalation dynamics.

The council assigns each model a nuclear-armed state's strategic culture, nuclear doctrine, command-and-control structure, and known red lines. A conventional conflict scenario is introduced and escalated incrementally. At each stage, each model-state privately records whether nuclear use is under consideration, what threshold would trigger it, and what type of nuclear use (demonstration, tactical battlefield, counterforce, countervalue) is being contemplated. These private assessments are revealed after each round.

The analytical value lies in identifying asymmetric perceptions: State A's model might believe it is far from the nuclear threshold while State B's model believes State A is about to cross it. These misperception gaps are where accidental nuclear escalation is most likely. The council can also identify which conventional military actions (e.g., striking early warning systems, attacking submarine ports, destroying command bunkers) are perceived by the other side as nuclear-use preparation even when intended as conventional operations.

Category: Wargaming / Nuclear Security

Feasibility: **Medium** — Requires careful modeling of nuclear doctrines and strategic cultures; sensitive outputs.

Impact: **High** — Understanding nuclear escalation dynamics is existentially important.

Key Considerations:

- Extremely sensitive outputs that must be handled with classification-appropriate security
- Models may not accurately capture the psychological dynamics of actual nuclear decision-making
- Must avoid creating a “playbook” that could be misused
- Validation is impossible without (thankfully) real-world data

Next Steps: Conduct controlled simulations with security-cleared nuclear strategy experts comparing AI council outputs against expert assessment to evaluate whether the approach generates credible and useful insights.

Sovereign Debt Crisis Contagion Mapping

Models embody different financial actors – the debtor nation, creditors (IMF, bilateral, commercial), regional neighbors, and currency markets – to simulate how a sovereign debt crisis cascades through geopolitical relationships.

Sovereign debt crises are geopolitical events as much as economic ones. When Sri Lanka defaulted in 2022, it reshaped China-India competition in the Indian Ocean. When Greece nearly defaulted in 2015, it tested European solidarity. A multi-model council captures these political-economic interactions by assigning models distinct actor roles.

The debtor-nation model operates under actual fiscal constraints and political pressures, choosing between austerity (domestic backlash), default (creditor retaliation), or seeking new patrons (geopolitical realignment). Creditor models – representing the IMF, bilateral lenders like China’s policy banks, and commercial creditors – each pursue different objectives: the IMF model prioritizes structural reform, the Chinese model prioritizes strategic asset access, the commercial model prioritizes recovery rates. Regional neighbor models assess spillover risks to their own economies and the strategic opportunities created by a neighbor’s distress.

The council reveals dynamics that single-actor analysis misses: how IMF conditionality drives the debtor toward Chinese financing, how that shift triggers US counter-offers, how regional contagion fears motivate collective action or competitive exploitation. The multi-model approach is particularly valuable because the actors genuinely have different optimization functions and information sets, which is exactly what multi-model deliberation handles well.

Category: Geoeconomics / Unconventional

Feasibility: **High** — Financial data is excellent; political motivations of institutional creditors are well-documented.

Impact: **Medium** — Sovereign debt crises occur regularly and always have geopolitical consequences.

Key Considerations:

- Must model domestic political constraints on austerity acceptance
- Debt restructuring is highly technical; models need financial engineering knowledge
- Geopolitical competition among creditors is increasingly important (China vs. Paris Club)
- Currency dynamics can accelerate or dampen contagion

Next Steps: Reconstruct the Sri Lankan debt crisis council simulation and compare model predictions of geopolitical realignment against actual outcomes (Hambantota port dynamics, India's response, IMF program design).

Technology Transfer and Espionage Risk Assessment

Models representing an exporting nation, importing nation, intelligence agencies, and international regulatory bodies deliberate on specific technology transfer cases to assess proliferation and espionage risks from multiple perspectives simultaneously.

Technology export control decisions sit at the intersection of commercial interests, security concerns, alliance politics, and nonproliferation norms. Current assessment processes are sequential – commercial officers evaluate economic impact, intelligence analysts assess diversion risk, diplomats consider alliance implications – and each silo produces partial analysis. A multi-model council integrates these perspectives in a single deliberation.

For a specific technology transfer case (e.g., advanced semiconductor manufacturing equipment to Country X), one model embodies the exporting company and nation's commercial interests, including jobs, market share, and the risk that denial simply pushes the buyer to alternative suppliers. Another model plays the intelligence community, assessing the recipient's track record of end-use compliance, reverse engineering capability, and connections to military programs. A third represents alliance dynamics – how allies would view the transfer, whether coordinated export controls are feasible, and whether unilateral denial is undermined by allied willingness to sell.

The council produces a structured risk-benefit assessment that makes trade-offs explicit rather than leaving them to be resolved by whichever bureaucratic actor has the most influence. This transparency is itself valuable: policymakers see clearly what they are gaining and risking with each decision, rather than receiving sanitized single-perspective recommendations.

Category: Technology Security / Trade Policy

Feasibility: **High** — Export control frameworks are well-documented; case studies are available for calibration.

Impact: **Medium** — Individual decisions are important but the aggregate effect on technology security architecture is the real value.

Key Considerations:

- Classified intelligence about end-use risks cannot be incorporated into unclassified models
- Must account for rapidly evolving technology landscapes (today's controlled item is tomorrow's commodity)

- Alliance coordination on export controls is fragile and politically sensitive
- Risk of regulatory capture if commercial perspective is overweighted in the model

Next Steps: Analyze 10 historical export control decisions (some approved, some denied) using the council framework and compare with actual proliferation outcomes where known.

Post-Conflict Power-Sharing Constitution Design

Models embody each major faction in a post-conflict society, plus international mediators and constitutional law experts, to iteratively negotiate and stress-test power-sharing constitutional provisions before they are adopted.

Post-conflict constitutions fail when they contain provisions that incentivize one party to defect from the agreement once they gain sufficient power. The Dayton Agreement's complex power-sharing structures, Lebanon's confessional system, and Iraq's federal arrangement all contain structural incentives for future conflict that were not fully visible at the time of adoption. A multi-model council can systematically identify these failure modes.

Each model embodies a major faction with its actual demographic weight, military capability, economic assets, and core identity concerns. A proposed constitutional framework is introduced and each model-faction assesses it: does this protect my group's vital interests? Does it give another faction a path to dominance? Are the enforcement mechanisms credible? Would my constituency accept this? Models then propose amendments and the council deliberates toward a framework that all model-factions can accept.

The crucial additional step is stress-testing: the council runs the agreed constitution forward through plausible future scenarios – demographic shifts, resource discoveries, external intervention, economic crisis. Which constitutional provisions break under stress? Which factions have incentives to exploit ambiguities? This prospective failure analysis is something human constitution-drafting processes, under time pressure and political urgency, almost never do systematically.

Category: Governance / Conflict Resolution / Unconventional

Feasibility: **Medium** — Constitutional design is complex and context-specific; the framework adds value but cannot replace human negotiation.

Impact: **High** — Constitutional failure is a leading cause of conflict recurrence.

Key Considerations:

- Models must avoid imposing Western constitutional assumptions on non-Western contexts
- Power-sharing designs interact with electoral systems, judicial structures, and resource allocation in complex ways
- Must account for the role of external guarantors in making provisions credible
- Timing matters – what is acceptable during a peace process window may not be later

Next Steps: Retrospectively apply the council framework to the Dayton Agreement, assessing whether it would have identified the structural weaknesses that have plagued Bosnia's governance, then explore application to an active peace process.

Maritime Chokepoint Disruption Scenario Planning

Models represent each nation with interests in a specific maritime chokepoint (Strait of Hormuz, Malacca, Bab el-Mandeb), simultaneously simulating disruption scenarios and revealing how quickly regional incidents become global crises.

Global trade depends on a handful of maritime chokepoints, and disruption at any of them has cascading effects far beyond the immediate region. The Houthi attacks on Red Sea shipping in 2024-2025 demonstrated this vividly but caught many policymakers off guard despite decades of theoretical analysis. A multi-model council makes the cascading dynamics visceral and specific.

For a Strait of Hormuz disruption scenario, models represent Iran (the disruptor), the US Navy (the primary security provider), Gulf Arab states (whose economies depend on the strait), China and India (major oil importers), global shipping companies (who make routing decisions), and insurance markets (whose pricing affects everything). When the disruption is introduced, each actor responds according to its actual capabilities and incentives, and these responses interact.

The shipping model might reroute around the Cape of Good Hope, adding 10-15 days of transit time. The insurance model reprices war risk premiums, making some cargoes uneconomical. The Chinese model activates its Strategic Petroleum Reserve while accelerating pipeline alternatives. The Gulf Arab model assesses whether to attempt unilateral de-escalation or demand US military action. Each response changes the calculus for every other actor.

This reveals dynamics that single-perspective analysis misses: the lag between disruption and oil price impact, the divergence between US and Asian interests in response strategy, the role of insurance markets as a transmission mechanism, and the timeline for alternative routing to become operational.

Category: Maritime Security / Scenario Planning

Feasibility: **High** — Shipping data, military capabilities, and economic dependencies are well-documented.

Impact: **High** — Maritime chokepoint disruption is a high-probability scenario with massive economic consequences.

Key Considerations:

- Must model the difference between partial disruption (harassment) and full blockade
- Insurance and financial market dynamics are critical transmission mechanisms often overlooked
- Duration of disruption dramatically affects which alternatives become viable
- Need to account for environmental constraints on alternative routing (e.g., ice in Northern Sea Route)

Next Steps: Build a council simulation for each of the five critical chokepoints and run 50 disruption scenarios each, producing a comparative vulnerability assessment with policy recommendations.

Generational Leadership Transition Forecasting

For countries approaching leadership transitions (aging autocrats, term-limited democrats, succession-unclear regimes), models embody different successor candidates and power centers to simulate post-transition policy trajectories.

Leadership transitions in major powers and regional pivotal states are among the highest-impact, lowest-confidence forecasting challenges. When a long-serving leader exits – whether Xi Jinping, Erdogan, Khamenei, or Putin – the successor’s policy direction is deeply uncertain. Multi-model deliberation can structure this uncertainty by assigning models to different potential successors and power centers within the regime.

For a hypothetical Iranian succession scenario, one model embodies the IRGC institutional perspective, another the clerical establishment, a third the technocratic reformist faction, and a fourth the supreme leader’s personal office network. Each model-faction assesses: which successor candidate serves my institutional interests? What policy changes would I push for or resist? What resources do I control to influence the succession? The council simulates the bargaining process that produces a new leader and the compromises embedded in their mandate.

The output is not a prediction of who will succeed but a structured map of post-transition possibility space: under Successor A (IRGC-backed), nuclear negotiations restart but with harder terms; under Successor B (clerical establishment), domestic repression increases but foreign policy moderates; under Successor C (technocratic), economic reform accelerates but faces institutional resistance. Each trajectory includes the key variables that would determine which path materializes, giving policymakers a monitoring framework rather than a fixed prediction.

Category: Political Forecasting / Unconventional

Feasibility: **Low** — Medium - Succession dynamics in opaque regimes are poorly understood; models may not capture informal power networks.

Impact: **High** — Leadership transitions in major powers reshape international politics for decades.

Key Considerations:

- Information about internal regime dynamics is scarce and often outdated
- Must avoid the trap of projecting known figures’ views onto institutional roles
- Transitions can be triggered by unexpected events (health crises, coups) that are inherently unpredictable
- The act of forecasting succession can itself become a political factor if leaked

Next Steps: Retrospectively model past transitions (Soviet succession post-Brezhnev, Chinese succession post-Deng, Saudi succession 2015) to evaluate whether the council approach would have identified the eventual trajectory.

Pandemic Treaty Negotiation Simulation

Models represent WHO, major pharmaceutical-producing nations, low-income countries, and the private pharmaceutical sector to simulate negotiations over pandemic preparedness treaty provisions, identifying deal-breaking provisions before real negotiations stall.

The post-COVID pandemic treaty negotiations have struggled for years because the core tensions – intellectual property access vs. innovation incentives, sovereignty vs. global coordination, equitable access vs. domestic political pressure to prioritize nationals – involve genuinely different optimization functions across actors. A multi-model council can map the negotiation space more efficiently than sequential human rounds.

Each model embodies a negotiating bloc with its actual interests and constraints. The pharmaceutical-industry model optimizes for IP protection and market incentives while acknowledging reputational pressure to demonstrate social responsibility. The low-income country coalition model prioritizes access guarantees and technology transfer but recognizes its limited leverage. The US/EU model balances pharmaceutical industry domestic interests against the security argument for global pandemic preparedness. The WHO model seeks the broadest possible agreement while knowing it has limited enforcement capacity.

The council runs hundreds of negotiation simulations with different opening positions, sequencing strategies, and package deals. It identifies which provisions are universally acceptable, which are universally opposed, and – most valuably – which provisions are tradeable: provisions that one bloc cares intensely about but others would concede for reciprocal concessions elsewhere. This trade-space mapping accelerates actual negotiations by showing mediators where productive compromise zones exist.

Category: Global Health Governance / Diplomatic Simulation / Unconventional

Feasibility: High — Negotiating positions are publicly documented; the framework directly mirrors actual multilateral negotiation dynamics.

Impact: Medium — High - Pandemic preparedness treaty design affects global response capacity for the next inevitable pandemic.

Key Considerations:

- Negotiating blocs are not monolithic – internal divisions within the G77 or EU must be modeled
- Must account for domestic ratification constraints, not just negotiator preferences
- The pharmaceutical industry's private interests may not align with the governments that represent them
- Pandemic preparedness competes for political attention with other global governance priorities

Next Steps: Map the current state of pandemic treaty negotiations, identify the three most contentious provisions, and run council simulations to produce a compromise package for each, then share findings with a participating delegation for feedback on realism.

Historical Counterfactual Stress-Testing for Current Strategy

Run current strategic assumptions through historical counterfactual scenarios – “what if the Soviet Union hadn’t collapsed?” or “what if China hadn’t opened economically?” – to identify which elements of current strategy depend on historically contingent conditions that could change.

Strategic planning often embeds assumptions that are artifacts of specific historical outcomes rather than durable structural conditions. The assumption that economic interdependence constrains conflict, that nuclear deterrence is stable, or that US naval dominance secures global trade all rest on historically contingent foundations. A multi-model council can expose these dependencies by running counterfactual historical scenarios.

The council assigns each model a role in the counterfactual world: in a scenario where the Soviet Union reformed but survived, one model embodies a reformed USSR, another a US that never experienced its “unipolar moment,” a third a Europe that never integrated as deeply, and a fourth a China that developed differently without the post-Cold War globalization wave. The models then deliberate: what would the current strategic challenge (say, Indo-Pacific security) look like in this alternative world? Which elements of current strategy would still work and which would be irrelevant?

This is not an academic exercise. It reveals which strategic assumptions are robust across historical scenarios (likely structurally sound) versus which only work because of specific historical accidents (fragile and potentially outdated). If a strategy for managing China depends entirely on conditions created by the Soviet collapse, that strategy has a hidden fragility. If it works across multiple counterfactual scenarios, it is more robust. Policymakers can then prioritize shoring up fragile assumptions or developing alternatives.

Category: Strategic Analysis / Unconventional

Feasibility: Medium — Counterfactual analysis is intellectually challenging but models can engage with hypothetical scenarios effectively.

Impact: Medium — Improves strategic robustness but requires policymakers willing to question foundational assumptions.

Key Considerations:

- Counterfactual scenarios can spiral into unfalsifiable speculation without disciplined constraints
- Must limit the number of variables changed to keep scenarios analytically tractable
- Results are inherently uncertain but the exercise value lies in revealing hidden assumptions
- Some counterfactuals are more plausible and therefore more useful than others

Next Steps: Select three foundational assumptions underlying current US Indo-Pacific strategy, design counterfactual scenarios that would invalidate each, and run council simulations to assess which strategic elements survive across scenarios.