

Use Cases for Karpathy's LLM Council Model

AI Ideation Run

Date: 2026-04-12

Ideas: 10

Depth: Standard

Medical Symptom Triage for Personal Health Decisions

Use a council of LLMs to evaluate personal health symptoms before deciding whether to seek urgent care, schedule a routine appointment, or self-manage.

When someone experiences ambiguous health symptoms — a persistent headache, unusual fatigue, or a new skin mark — they often consult a single AI or web search and get either alarmist or dismissive responses depending on which source they hit. The council approach is uniquely suited here because different LLMs have different training biases around medical caution: some over-triage, others under-triage.

By querying multiple models and looking at the distribution of responses — how many say “see a doctor urgently,” how many say “monitor for a week,” how many say “likely benign” — the user gets a much more calibrated sense of the actual risk landscape. The consensus signal is far more valuable than any single model’s opinion. If 4 out of 5 models flag something as worth investigating, that carries real weight. If only 1 out of 5 does, the user can weigh that appropriately.

This isn’t a replacement for medical advice, but it’s a significant upgrade over the current pattern of Googling symptoms and landing on whichever result confirms existing anxiety or complacency.

Category: Personal Health

Feasibility: High — All major LLMs already handle medical symptom queries; the orchestration layer is straightforward.

Impact: High — Health decisions are high-stakes and highly susceptible to single-source bias.

Key Considerations:

- Must include clear disclaimers that this is not medical diagnosis
- Models may share training data biases (e.g., all trained on the same medical corpora), reducing true diversity
- Liability questions if users rely on council consensus to delay care
- Could integrate confidence scores alongside the consensus vote

Next Steps: Build a prototype that queries 3-5 models with structured symptom descriptions, classifies each response into triage categories, and presents the distribution to the user alongside the raw reasoning from each model.

Enterprise Contract Clause Risk Assessment

Run contract clauses through a council of LLMs to identify risks, unusual terms, and negotiation leverage points that any single model might miss.

Legal contract review is one of the clearest cases where multiple perspectives genuinely surface different findings. Different LLMs have been trained on different legal corpora and embed different heuristics about what constitutes “standard” vs. “unusual” in contract language. A clause that one model flags as aggressive might be treated as boilerplate by another — and that disagreement itself is informative.

The council approach lets a legal team run key clauses through multiple models and focus human review on the areas of highest model disagreement. High consensus that a clause is standard means it can be fast-tracked. High disagreement signals genuine ambiguity worth escalating to senior counsel. This prioritization alone can save significant review hours on large deals.

The synthesis step is particularly powerful here: rather than just voting, the council can produce a merged analysis that captures the union of all risks identified across models, ensuring nothing falls through the cracks of any single model's blind spots.

Category: Business Legal

Feasibility: High — Contract review is already a major LLM use case; adding the council layer is incremental.

Impact: High — Missed contract risks can cost millions; the marginal cost of querying additional models is trivial by comparison.

Key Considerations:

- Confidentiality: enterprise contracts may contain sensitive terms, requiring on-premise or private API deployments
- Need to normalize different models' risk rating scales before aggregation
- Most valuable for non-routine contracts (M&A, partnerships) rather than high-volume standard agreements
- Could pair with human-in-the-loop where disagreement exceeds a threshold

Next Steps: Pilot with a corporate legal team on a set of recently negotiated contracts (where outcomes are known), measuring whether council disagreement correlates with clauses that actually required negotiation changes.

Personal Life Decision Sounding Board

Use the council as a structured decision-making tool for major personal life choices — career moves, relocations, relationship decisions — where the diversity of “advisor perspectives” is the core value.

People facing major life decisions often seek advice from multiple friends, family members, or mentors precisely because different people bring different frameworks. The LLM council formalizes this: each model becomes a different “advisor” that may weigh factors differently. One model might emphasize financial security, another personal fulfillment, another risk tolerance, another family impact.

What makes this particularly suited to the council approach is that the user doesn't want a single “correct” answer — they want to see the landscape of reasonable perspectives. The council's aggregation can surface which factors most advisors agree are important (high consensus) and which are genuinely trade-offs where reasonable people disagree (high variance). This reframes the decision from “what should I do” to “what are the real tensions I need to resolve.”

The approach can be enhanced by prompting each model with a different advisory persona (the pragmatist, the optimist, the devil's advocate) to ensure genuine perspective diversity rather than relying on incidental training differences alone.

Category: Personal Decision-Making

Feasibility: **High** — No special infrastructure needed; can be implemented with simple prompt orchestration.

Impact: **Medium** — Genuinely helpful for the subset of people who think through decisions systematically, but won't replace emotional processing or trusted human advisors.

Key Considerations:

- Risk of decision paralysis if the council surfaces too many conflicting perspectives without clear weighting
- Works best when the user has already articulated the decision and key factors clearly
- Could incorporate the user's stated values as weighting criteria for the final synthesis
- Privacy considerations for deeply personal decisions being sent to multiple API providers

Next Steps: Design a decision-framing template that structures the input (decision, options, factors, values) and an output format that clearly separates consensus factors from genuine tensions.

Code Review and Architecture Decision Records

Submit code changes or architecture proposals to a council of LLMs for review, using disagreement patterns to identify the genuinely contentious design decisions worth human debate.

In software engineering, code review quality varies enormously based on the reviewer's experience and biases. Some reviewers obsess over style, others focus on performance, others on security, others on maintainability. The LLM council can simulate this diversity systematically: each model brings slightly different training biases about what constitutes "good" code.

The high-value output isn't the individual reviews — it's the disagreement map. When all models agree something is problematic, it's probably a real issue. When all agree something is fine, it probably is. The interesting cases are where models disagree: one flags a function as too complex while others see it as appropriately detailed, or one suggests a different architectural pattern while others endorse the current approach. These disagreement points map directly to the design tensions that benefit most from human judgment.

For Architecture Decision Records (ADRs), this is even more powerful. Presenting a proposed architecture to the council and asking each to argue for or against it — then synthesizing the arguments — produces a much richer decision record than any single model's analysis.

Category: Business Software Engineering

Feasibility: **High** — Integrates naturally into existing CI/CD pipelines and PR workflows.

Impact: **Medium** — Most valuable for senior engineering decisions rather than routine code review.

Key Considerations:

- Models may converge on the same “best practices” from shared training data, reducing genuine diversity
- Need to distinguish style disagreements (low value) from architectural disagreements (high value)
- Cost scales with code change size and number of models; may want to reserve council review for significant changes
- Could integrate with existing code review tools (GitHub, GitLab) as a pre-review step

Next Steps: Implement as a GitHub Action that triggers council review on PRs tagged with a specific label, producing a structured disagreement report as a PR comment.

Investment Thesis Stress Testing

Present an investment thesis to a council of LLMs and use the distribution of bull/bear assessments as a signal for how robust the thesis is.

Investment decisions are plagued by confirmation bias — once someone develops a thesis, they tend to seek confirming evidence. A single LLM asked to evaluate a thesis will often mirror the framing of the prompt. But a council of models, especially when some are explicitly prompted to find weaknesses, produces a genuine diversity of analytical perspectives.

The council’s consensus pattern becomes a rough proxy for thesis robustness. If the user presents “I think Company X will benefit from the AI infrastructure buildout because...” and 4 out of 5 models agree with the core logic but identify different risk factors, that’s a reasonably strong thesis with known risks. If 3 out of 5 models fundamentally disagree with the premise, the thesis deserves much deeper scrutiny before capital allocation.

This works for both personal portfolio decisions and professional investment analysis. The key insight is that the council doesn’t predict markets — it stress-tests reasoning. High model agreement on a thesis doesn’t mean it’s correct, but high disagreement reliably signals that the thesis rests on assumptions that are genuinely debatable.

Category: Business/Personal Finance

Feasibility: **High** — Financial analysis is well within current LLM capabilities; aggregation is the novel layer.

Impact: **High** — Even marginal improvements in investment decision quality have outsized financial returns.

Key Considerations:

- Models may share biases from similar financial training data (e.g., all bullish on certain sectors)
- Works best for thesis stress-testing, not for generating novel investment ideas
- Should not be used as the sole basis for investment decisions — it’s an analytical tool, not a prediction engine
- Regulatory considerations if marketed as financial advice

Next Steps: Back-test by running historical investment theses through the council and correlating council consensus levels with actual outcomes over 1-3 year horizons.

Content Moderation Calibration for Online Platforms

Use a council of LLMs to calibrate content moderation decisions, particularly for edge cases where the line between acceptable and policy-violating content is genuinely ambiguous.

Content moderation at scale is one of the hardest problems in platform governance. The hardest cases aren't the clear violations — they're the edge cases where reasonable people disagree about whether something crosses the line. Satire vs. hate speech. Graphic journalism vs. gratuitous violence. Cultural context that makes something offensive in one context and normal in another.

The council approach is architecturally ideal for this because the distribution of model judgments directly maps to the ambiguity of the content. If 5 out of 5 models flag something, auto-action is safe. If 0 out of 5 flag it, auto-approve is safe. The 2-out-of-5 and 3-out-of-5 cases are precisely the ones that should be routed to human reviewers — and the council's reasoning provides those reviewers with a pre-structured analysis of why the content is ambiguous.

This approach also provides a natural audit trail and consistency mechanism. Platforms can track whether their council's consensus patterns align with human reviewer decisions over time and recalibrate accordingly.

Category: Business Platform Operations

Feasibility: **Medium** — Requires careful prompt engineering for each policy category and ongoing calibration against human decisions.

Impact: **High** — Content moderation errors (both false positives and false negatives) have massive reputational and legal consequences for platforms.

Key Considerations:

- Models may share systematic biases (e.g., all overly cautious about certain topics due to RLHF training)
- Latency requirements for real-time moderation may conflict with querying multiple models
- Cultural and linguistic context varies enormously; council may need region-specific configuration
- Must avoid using the council as a shield against accountability for moderation decisions

Next Steps: Partner with a mid-size platform to run the council in shadow mode (parallel to existing moderation) for 3 months, measuring agreement rates with human reviewers and identifying systematic blind spots.

Personal Writing Style and Tone Calibration

Before sending an important email, message, or document, run it through a council to assess how the tone will be perceived — is it too aggressive, too passive, appropriately assertive, or potentially ambiguous?

Tone misreading is one of the most common sources of interpersonal friction in written communication. The writer's intended tone and the reader's perceived tone often diverge, especially in

professional contexts, cross-cultural communication, or emotionally charged situations. A single LLM asked “how does this sound?” gives one interpretation. A council gives a distribution.

If you’re about to send a difficult email to a colleague and 4 out of 5 models read it as “direct but professional” while 1 reads it as “passive-aggressive,” that’s useful calibration — most readers will probably receive it well, but there’s a non-trivial risk of misreading. If 3 out of 5 read it as “confrontational,” you probably want to revise before sending.

The council approach is specifically superior to single-model tone checking because tone perception is inherently subjective and distribution-shaped. There is no single “correct” reading — the value is in seeing the range of plausible interpretations.

Category: Personal Communication

Feasibility: **High** — Simple to implement; tone assessment is a well-understood LLM capability.

Impact: **Medium** — High value for specific high-stakes messages (performance reviews, negotiations, conflict resolution), lower value for routine communication.

Key Considerations:

- Could extend to cultural tone calibration (how would this read in Japanese business culture vs. American vs. Israeli?)
- Privacy implications of routing sensitive communications through multiple API providers
- Works best when the user specifies the intended tone so the council can assess alignment
- Could provide specific revision suggestions weighted by the most common misreadings

Next Steps: Build a browser extension or email client plugin that offers one-click council tone assessment before sending, with a simple visualization of the interpretation distribution.

Regulatory Compliance Interpretation for Multi-Jurisdiction Businesses

Use a council to interpret how ambiguous regulatory language applies to a specific business practice, surfacing the range of plausible interpretations rather than a single (potentially over-confident) reading.

Businesses operating across jurisdictions constantly face regulatory language that is deliberately or incidentally ambiguous. “Reasonable security measures,” “adequate notice,” “material change” — these terms mean different things to different regulators, courts, and legal scholars. A single LLM asked “does our practice comply with GDPR Article X?” will give a confident answer that may not represent the full range of legitimate interpretations.

The council approach is ideal because regulatory interpretation is fundamentally about the distribution of plausible readings. If all models agree a practice is clearly compliant or clearly non-compliant, the compliance team can act with confidence. If the council is split, that’s a genuine signal that the regulation is ambiguous on this point and the business should seek formal legal counsel or adopt a more conservative interpretation.

This is particularly powerful for emerging regulations (AI governance, data sovereignty, ESG reporting) where there's limited case law and the "correct" interpretation is genuinely unsettled. The council's disagreement pattern maps directly to regulatory risk.

Category: Business Legal/Compliance

Feasibility: **Medium** — Requires models with strong legal training and careful prompt construction to specify jurisdiction and context.

Impact: **High** — Regulatory missteps can result in massive fines and operational disruption; better interpretation reduces risk significantly.

Key Considerations:

- Models may lack jurisdiction-specific training for smaller or newer regulatory regimes
- Must not be used as a substitute for formal legal advice, but as a screening and prioritization tool
- Particularly valuable for "right to be forgotten" requests, data transfer assessments, and AI use disclosures
- Could maintain a rolling knowledge base of past interpretations to track how consensus evolves as regulations mature

Next Steps: Pilot with a compliance team handling GDPR and AI Act obligations, running ambiguous requirements through the council and comparing consensus levels with outside counsel's actual assessments.

Hiring Decision Bias Reduction

Run anonymized candidate profiles through a council of LLMs for initial assessment, using consensus and disagreement patterns to identify where evaluator bias (human or AI) is most likely influencing decisions.

Hiring is one of the highest-stakes decisions organizations make, and it's riddled with both conscious and unconscious bias. A single AI screening tool embeds the biases of its training data. A single human reviewer embeds their personal biases. The council approach doesn't eliminate bias — but it makes bias detectable by surfacing disagreement.

If a council of 5 models is asked to assess a candidate's fit for a role and 4 agree the candidate is strong while 1 rates them weak, the interesting question is: what is that 1 model seeing that the others aren't? Is it picking up on a legitimate concern, or is it exhibiting a bias that the others don't share? By making these disagreements visible and traceable to specific candidate attributes, hiring teams can audit their process for systematic patterns.

The approach is especially powerful when combined with candidate anonymization. If the council's consensus changes significantly when demographic information is added vs. removed, that's a direct measure of bias in the models — and a useful calibration tool for the human process as well.

Category: Business Human Resources

Feasibility: **Medium** — Requires careful anonymization, structured candidate profiles, and integration with existing ATS systems.

Impact: High — Better hiring decisions compound over years; bias reduction has both ethical and business performance benefits.

Key Considerations:

- All models may share systematic biases (e.g., preference for prestigious university names) — council diversity doesn't guarantee bias elimination
- Must comply with emerging AI hiring regulations (EU AI Act, NYC Local Law 144)
- Works best as a supplement to human decision-making, not a replacement
- Need to validate that model disagreement actually correlates with meaningful assessment differences rather than noise

Next Steps: Design a controlled experiment: run 100 anonymized candidate profiles through a 5-model council, then compare council consensus with actual hiring outcomes and post-hire performance ratings.

Geopolitical Scenario Planning and Wargaming

Use a council of LLMs to generate and stress-test geopolitical scenarios, where each model's different training biases and analytical frameworks produce genuinely diverse forecasts that single-model analysis misses.

Geopolitical analysis is perhaps the domain where the council approach has the strongest theoretical justification. Geopolitical forecasting is notoriously difficult precisely because outcomes depend on the interaction of many actors with different motivations, information, and decision-making frameworks. A single analyst (human or AI) inevitably brings a dominant analytical lens — realist, liberal institutionalist, constructivist — that shapes which scenarios they consider plausible.

A council of LLMs, especially when prompted with different analytical frameworks or assigned different national perspectives, produces a genuinely broader scenario space. Model A might emphasize economic interdependence as a constraint on conflict, while Model B emphasizes domestic political pressures as an escalation driver. The distribution of scenarios across the council — which escalation paths appear in most models' analyses, which appear in only one — is itself the most valuable output.

For businesses with geopolitical exposure (supply chains, market access, regulatory risk) and for policy analysts, this approach provides structured scenario diversity that's extremely difficult to achieve through other means. The council's consensus scenarios represent the "base case" while outlier scenarios represent the tail risks worth hedging against.

Category: Business Strategy / Policy Analysis

Feasibility: Medium — Requires sophisticated prompt engineering to elicit genuinely different analytical frameworks rather than superficially different phrasings.

Impact: High — Geopolitical surprises (wars, sanctions, regime changes) can be catastrophic for unprepared organizations; broader scenario coverage improves preparedness.

Key Considerations:

- Models may converge on Western-centric analytical frameworks due to English-language training data dominance
- Deliberate prompting with specific analytical lenses (assign each model a different IR theory or national perspective) is essential for true diversity
- Most valuable for medium-term scenario planning (6-24 months) rather than short-term event prediction
- Could integrate with structured forecasting methodologies (Superforecasting, scenario matrices)

Next Steps: Run a retrospective test: apply the council to a geopolitical situation from 2 years ago, generate scenarios, and compare the council's scenario distribution with what actually happened — measuring whether the council would have included the actual outcome in its scenario set.

Generated by Claude (Anthropic) as part of a structured AI ideation run.